

REPORT

- Sauatbek Zhanserik 220103332 15-P
- Aktureyev Orazali 220103173 17-P
- Zharylkassynov Galymbek 220103251 15-P
- Abzalbek Mukhamedali 220103209 18-P
- Kantore Arman 220103282 19-P

1. Collecting the data

We decided to parse the kolesa.kz site to find some information related to cars in Kazakhstan. So I started the parsing stage.

First, I imported all the necessary libraries

```
import requests
from bs4 import BeautifulSoup
import pandas as pd
```

```
base_url = "https://kolesa.kz/cars/?auto-car-grbody="
```

base_url: This is the base URL of the website being scraped, which displays car listings.

```
headers = {
    "User-Agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/91.0.4472.124 Safari/537.36"
}
```

headers: This simulates a browser visit by adding a User-Agent. It prevents the server from blocking requests that appear to come from scripts.

```
all_car_data = []
```

all_car_data: An empty list to store the scraped car data.

```
max_pages = 100  
page = 1
```

max_pages: The maximum number of pages to scrape.

page: A counter to keep track of the current page being scraped.

```
while page <= max_pages:
```

This loop iterates over pages of car listings until the specified maximum (max_pages) is reached or there are no more listings.

```
url = f"{base_url}&page={page}"
```

url: Constructs the URL for the current page by appending the page number to the base_url.

```
response = requests.get(url, headers=headers)  
response.raise_for_status()  
print(f"Page {page} fetched successfully.")
```

requests.get: Sends a GET request to the URL.

raise_for_status: Throws an exception if the request fails.

```
soup = BeautifulSoup(response.text, "html.parser")
```

BeautifulSoup: Parses the HTML content of the webpage, making it easier to navigate and extract data.

```
car_listings = soup.find_all("div", class_="a-list__item")

if not car_listings:
    print("No more listings found. Ending scraping.")
    break
```

find_all: Searches for all div elements with the class a-list__item, which likely represents individual car listings on the page.

```
for car in car_listings:
    try:

        title = car.find("h5", class_="a-card__title").get_text(strip=True)

        price = car.find("span", class_="a-card__price").get_text(strip=True)

        details = car.find("p", class_="a-card__description").get_text(strip=True)

        location = car.find("span", class_="a-card__param").get_text(strip=True)
```

find: Locates specific elements within each car listing.

- **title:** The car's title or model, found in an <h5> tag with the class a-card__title.
- **price:** The car's price, found in a tag with the class a-card__price.
- **details:** Additional details about the car, found in a <p> tag with the class a-card__description.
- **location:** The car's location, found in a tag with the class a-card__param.

```

        all_car_data.append({
            "Title": title,
            "Price": price,
            "Details": details,
            "Location": location
        })
    except AttributeError:
        continue

```

The extracted details are stored as a dictionary and appended to the `all_car_data` list.

```

page += 1

```

Moves to the next page for scraping.

```

df = pd.DataFrame(all_car_data)

```

pd.DataFrame: Converts the list of dictionaries (`all_car_data`) into a structured DataFrame, making the data easier to analyze or export.

```

df.to_csv('test.csv')

```

```

df = pd.read_csv(r"C:\Users\Admin\test.csv")
df.head()

```

	Unnamed: 0	Title	Price	Details	Location
0	0	BA3 (Lada) Lada 2121	5 950 000₸	2020 г., Б/у внедорожник, 1.7 л, бензин, КПП м...	Семей
1	1	Subaru Legacy	6 500 000₸	2004 г., Б/у седан, 2 л, бензин, Правый руль, ...	Алматы
2	2	Mercedes-Benz C 180	2 200 000₸	1996 г., Б/у седан, 1.8 л, бензин, КПП механик...	Жанаозен
3	3	Genesis G70	23 500 000₸	2021 г., Б/у седан, 3.3 л, бензин, КПП автомат...	Алматы
4	4	Hyundai Santa Fe	14 790 152₸	2021 г., Б/у кроссовер, 2.5 л, бензин, КПП авт...	Павлодар

Because my parsed dataset was not perfect I saved it like test.csv to do a future data preprocessing stage.

2. Data preprocessing

```
df[['Brand', 'Model']] = df['Title'].str.extract(r'^(?P<Brand>\S+)\s+(?P<Model>.+)$')
```

I divided the Title column to Brand and Model columns for future visualizations part,

```
details_split = df['Details'].str.extract(
    r'^(?P<Year>\d{4}) r\., (?P<Condition>Б/у|Новые)? (?P<Type>[^,]+), (?P<Engine_Size>\d+\.\d+ л), (?P<Fuel_Type>[^,]+), (?P<Transmission>КПП [^,]+), c
)

df = pd.concat([df, details_split], axis=1)

df.head()
```

There I also divided the details column into several other columns with features and info about the cars and conceited it with my dataframe.

```
df.isna().sum()

Unnamed: 0      0
Title           0
Price           0
Details         0
Location        0
Brand           0
Model           0
Year           1048
Condition       1048
Type            1048
Engine_Size     1048
Fuel_Type       1048
Transmission    1048
Mileage         1048
Color           1048
Additional      1048
dtype: int64
```

```
df = df.dropna()
```

I saw a lot of missing values and I decided to drop them all because it would be very problematic to fill them with unknown.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 952 entries, 0 to 1999  
Data columns (total 16 columns):  
#   Column                Non-Null Count  Dtype  
---  ---  
0   Unnamed: 0             952 non-null   int64  
1   Title                  952 non-null   object  
2   Price                  952 non-null   object  
3   Details                952 non-null   object  
4   Location               952 non-null   object  
5   Brand                  952 non-null   object  
6   Model                  952 non-null   object  
7   Year                   952 non-null   object  
8   Condition              952 non-null   object  
9   Type                   952 non-null   object  
10  Engine_Size            952 non-null   object  
11  Fuel_Type              952 non-null   object  
12  Transmission           952 non-null   object  
13  Mileage                952 non-null   object  
14  Color                  952 non-null   object  
15  Additional             952 non-null   object  
dtypes: int64(1), object(15)  
memory usage: 126.4+ KB
```

There you see that all the types of columns are in the object type, which means I should change the types where it's necessary.

```
df = df.drop('Details', axis=1)
```

```
df = df.drop('Unnamed: 0', axis=1)
```

There I dropped unnecessary columns

df.head()

	Title	Price	Location	Brand	Model	Year	Condition	Type	Engine_Size	Fuel_Type	Transmission	Mileage	Color	Additional
0	BA3 (Lada) Lada 2121	5 950 000Т	Семей	BA3	(Lada) Lada 2121	2020	Б/у	внедорожник	1.7 л	бензин	КПП механика	29 000	белый	литые диски, тонировка, ветровики,...
3	Genesis G70	23 500 000Т	Алматы	Genesis	G70	2021	Б/у	седан	3.3 л	бензин	КПП автомат	16 500	белый	металлик
4	Hyundai Santa Fe	14 790 152Т	Павлодар	Hyundai	Santa Fe	2021	Б/у	кроссовер	2.5 л	бензин	КПП автомат	107 000	белый	литые диски, ксенон, кожа, аудиосист...
7	Ravon Nexia R3	5 200 000Т	Кызылорда	Ravon	Nexia R3	2020	Б/у	седан	1.5 л	бензин	КПП автомат	118 000	серый	комбинированный, климат-контроль, Машина...
11	Toyota Camry	18 000 000Т	Астана	Toyota	Camry	2020	Б/у	седан	2.5 л	бензин	КПП автомат	39 000	белый	металлик, кожа, кондиционер, Вся родная и...

That's how the dataset looks so far.

```
desired_order = [
    'Title', 'Brand', 'Model', 'Year', 'Price', 'Condition', 'Type', 'Engine_Size', 'Fuel_Type',
    'Transmission', 'Mileage', 'Color', 'Additional'
]

df = df[desired_order]
```

I decided to change the order of the columns for better readability.

```
df['Price'] = df['Price'].str.replace(r'\D', '', regex=True).astype(int)
```

There I extracted the value of the price column and change the data type into an integer.

```
df['Engine_Size'] = df['Engine_Size'].str.extract(r'(\d+\.\d+|\d+)').astype(float)
```

I did the same to the engine_size column.

```
df['Mileage'] = df['Mileage'].str.replace(' ', '').astype(int)
```

and for the mileage column.
This is the final dataset

	Title	Brand	Model	Year	Price	Condition	Type	Engine_Size	Fuel_Type	Transmission	Mileage	Color	Additional
0	BA3 (Lada) Lada 2121	BA3	(Lada) Lada 2121	2020	5950000	Б/у	внедорожник	1.7	бензин	КПП механика	29000	белый	литые диски, тонировка, ветровики,...
3	Genesis G70	Genesis	G70	2021	23500000	Б/у	седан	3.3	бензин	КПП автомат	16500	белый	металлик
4	Hyundai Santa Fe	Hyundai	Santa Fe	2021	14790152	Б/у	кроссовер	2.5	бензин	КПП автомат	107000	белый	литые диски, ксенон, кожа, аудиосист...
7	Ravon Nexia R3	Ravon	Nexia R3	2020	5200000	Б/у	седан	1.5	бензин	КПП автомат	118000	серый	комбинированный, климат-контроль, Машина...
11	Toyota Camry	Toyota	Camry	2020	18000000	Б/у	седан	2.5	бензин	КПП автомат	39000	белый	металлик, кожа, кондиционер, Вся родная и...
12	Volkswagen Golf	Volkswagen	Golf	1998	2550000	Б/у	хэтчбек	1.8	бензин	КПП автомат	241524	белый	металлик, литые диски, тонировка, спой...
14	BA3 (Lada) Priora 2170	BA3	(Lada) Priora 2170	2015	3950000	Б/у	седан	1.6	бензин	КПП механика	174547	серебристый	металлик, литые диски, тонировка,...
16	Toyota Highlander	Toyota	Highlander	2017	19500000	Б/у	кроссовер	3.5	бензин	КПП автомат	117000	черный	металлик, литые диски, тонировка, л...

```
df.to_csv('python_final2.csv')
```

Then I saved the final dataset to a .csv

```
df = df.drop('Unnamed: 0', axis=1)
```

```
df.head()
```

	Title	Brand	Model	Year	Price	Condition	Type	Engine_Size	Fuel_Type	Transmission	Mileage	Color	Additional
0	BA3 (Lada) Lada 2121	BA3	(Lada) Lada 2121	2020	5950000	Б/у	внедорожник	1.7	бензин	КПП механика	29000	белый	литые диски, тонировка, ветровики,...
1	Genesis G70	Genesis	G70	2021	23500000	Б/у	седан	3.3	бензин	КПП автомат	16500	белый	металлик
2	Hyundai Santa Fe	Hyundai	Santa Fe	2021	14790152	Б/у	кроссовер	2.5	бензин	КПП автомат	107000	белый	литые диски, ксенон, кожа, аудиосист...
3	Ravon Nexia R3	Ravon	Nexia R3	2020	5200000	Б/у	седан	1.5	бензин	КПП автомат	118000	серый	комбинированный, климат-контроль, Машина...
4	Toyota Camry	Toyota	Camry	2020	18000000	Б/у	седан	2.5	бензин	КПП автомат	39000	белый	металлик, кожа, кондиционер, Вся родная и...

Then we started to analyze our dataset.

```
df.describe
```

<bound method NDFrame.describe of								Title	Brand
0	BA3 (Lada)	Lada 2121	BA3	(Lada) Lada 2121	2020	5950000			
1	Genesis G70	Genesis	G70	2021	23500000				
2	Hyundai Santa Fe	Hyundai	Santa Fe	2021	14790152				
3	Ravon Nexia R3	Ravon	Nexia R3	2020	5200000				
4	Toyota Camry	Toyota	Camry	2020	18000000				
..				
947	BA3 (Lada) Granta 2190	BA3	(Lada) Granta 2190	2013	2870000				
948	Kia Sportage	Kia	Sportage	2023	18000000				
949	Hyundai Santa Fe	Hyundai	Santa Fe	2023	19700000				
950	Kia Sorento	Kia	Sorento	2021	19000000				
951	BA3 (Lada) 2112	BA3	(Lada) 2112	2000	650000				

	Condition	Type	Engine_Size	Fuel_Type	Transmission	Mileage
0	Б/у	внедорожник	1.7	бензин	КПП механика	29000
1	Б/у	седан	3.3	бензин	КПП автомат	16500
2	Б/у	кроссовер	2.5	бензин	КПП автомат	107000
3	Б/у	седан	1.5	бензин	КПП автомат	118000
4	Б/у	седан	2.5	бензин	КПП автомат	39000
..
947	Б/у	седан	1.6	бензин	КПП механика	157000
948	Б/у	кроссовер	2.5	бензин	КПП автомат	8000
949	Б/у	кроссовер	2.5	бензин	КПП автомат	20000
950	Б/у	кроссовер	2.5	бензин	КПП автомат	40000
951	Б/у	хэтчбек	1.6	бензин	КПП механика	222222


```
df.dtypes
```

Title	object
Brand	object
Model	object
Year	int64
Price	int64
Condition	object
Type	object
Engine_Size	float64
Fuel_Type	object
Transmission	object
Mileage	int64
Color	object
Additional	object
dtype:	object

```
df.info()
```

```
df.isna().sum()
```

#	Column	Non-Null Count	Dtype
0	Title	952 non-null	object
1	Brand	952 non-null	object
2	Model	952 non-null	object
3	Year	952 non-null	int64
4	Price	952 non-null	int64
5	Condition	952 non-null	object
6	Type	952 non-null	object
7	Engine_Size	952 non-null	float64
8	Fuel_Type	952 non-null	object
9	Transmission	952 non-null	object
10	Mileage	952 non-null	int64
11	Color	952 non-null	object
12	Additional	952 non-null	object

Title	0
Brand	0
Model	0
Year	0
Price	0
Condition	0
Type	0
Engine_Size	0
Fuel_Type	0
Transmission	0
Mileage	0
Color	0
Additional	0
dtype:	int64

As you can see here, we used some functions that will give us the descriptive analysis of our dataset.

```
df_price_mean = df['Price'].mean()
df_engine_mean = df['Engine_Size'].mean()
df_mileage_mean = df['Mileage'].mean()

print('Price mean: ', df_price_mean)
print('Engine size mean: ', df_engine_mean)
print('Mileage mean: ', df_mileage_mean)

Price mean:  10518212.87184874
Engine size mean:  2.521323529411765
Mileage mean:  156474.58088235295
```

I used the mean() function to identify average values for integer-type columns.

```
df_price_max = df['Price'].max()
df_engine_max = df['Engine_Size'].max()
df_mileage_max = df['Mileage'].max()

df_price_min = df['Price'].min()
df_engine_min = df['Engine_Size'].min()
df_mileage_min = df['Mileage'].min()

print('Price minimum and maximum: ', df_price_min, ' ', df_price_max)
print('Engine size minimum and maximum: ', df_engine_min, ' ', df_engine_max)
print('Mileage minimum and maximum: ', df_mileage_min, ' ', df_mileage_max)

Price minimum and maximum:  300000  80000000
Engine size minimum and maximum:  0.1  6.6
Mileage minimum and maximum:  1  888888
```

Also, I used max() and min() to identify minimum and maximum values on each column.

```
skewness = df['Year'].skew()
print(f"Skewness of Year: {skewness}")

kurtosis = df['Year'].kurt()
print(f"Kurtosis of Year: {kurtosis}")

Skewness of Year: -0.8853670970001123
Kurtosis of Year: 0.29032769010037374
```

Then I tried to find skewness and kurtosis for better analysis in the future. Our skewness is the negative. It means that most of our data in the right side of our normal distribution.

```
from sklearn.preprocessing import MinMaxScaler

normalizer = MinMaxScaler()


normalized_data = normalizer.fit_transform(df[['Year', 'Price', 'Engine_Size', 'Mileage']])
normalized_df = pd.DataFrame(normalized_data, columns=['Year', 'Price', 'Engine_Size', 'Mileage'])

print("Normalized Data:")
print(normalized_df)
```

Normalized Data:

	Year	Price	Engine_Size	Mileage
0	0.913043	0.070891	0.246154	0.032624
1	0.934783	0.291092	0.492308	0.018561
2	0.934783	0.181809	0.369231	0.120374
3	0.913043	0.061481	0.215385	0.132749
4	0.913043	0.222083	0.369231	0.043874
..
947	0.760870	0.032246	0.230769	0.176624
948	0.978261	0.222083	0.369231	0.008999
949	0.978261	0.243413	0.369231	0.022499
950	0.934783	0.234630	0.369231	0.044999
951	0.478261	0.004391	0.230769	0.249999

Then we used MinMaxScaler() to normalize our data for use in the future.



This is our central tendency. From this figure, we can observe that our mean and mode are in one position.

```
grouped_stats = df.groupby('Year').agg({'Price': ['mean', 'sum', 'std']})
grouped_stats
```

We used this aggregation function to find the most popular year of cars, which are on the website.

```
grouped_stats_name = df.groupby('Title').agg({'Price': ['mean', 'sum']})
sorted_df = grouped_stats_name.sort_values(by=('Price', 'sum'), ascending=False)
sorted_df.head()
```

Same thing but with brands.

```
unsorted = df.groupby('Title')['Mileage'].max().reset_index()
sorted = unsorted.sort_values(by='Mileage', ascending=True)
sorted
```

With this function, we can find cars that can pass a lot of kilometers. It will help if you are looking for a car.

```
mileage_range_type = df.groupby('Type')['Mileage'].agg(['min', 'max']).reset_index()
mileage_range_type
```


Here we can identify the best type of vehicle, for passing a lot of kilometers. It is a necessary part of cars.

```
average_price_by_brand = df.groupby('Brand')['Price'].mean()
sorted = average_price_by_brand.sort_values(ascending=False)
sorted
```

This is the distribution that shows the price differences between the brands.

```
import plotly.express as px


fig = px.histogram(df, x='Price', nbins=30, title='Price Distribution',
                  labels={'Price': 'Vehicle Price'},
                  color_discrete_sequence=['blue'])
fig.update_layout(bargap=0.2)
fig.show()
```



We used the “Plotly” library for our visualizations. Because It is the interactive library of Python.


Here you can see the price distribution for all cars.

```
fig = px.scatter(df, x='Year', y='Price', title='Price vs. Year',
                labels={'Year': 'Manufacturing Year', 'Price': 'Vehicle Price'},
                color='Type',
                size='Engine_Size',
                hover_data=['Brand', 'Model'])
fig.update_traces(marker=dict(opacity=0.7))
fig.show()
```




This is the price, year, and type of the car on the scatter plot.


Brand-Wise Price Comparison



Here you can see price distribution by brands. But you also can see the price differences between some models of the types.



Here you can find the most popular fuel types. (without oil)



This is the price for the any type of the fuel.

Data Aggregation and Grouping Operations

My part (Daniyar) here is started with Data Aggregation and Grouping, so I will use such functions as: **groupby()**, **pivot()**, **unstacked()**.

groupby(): This is used to group data by one or more columns, then we can apply aggregation functions like **sum**, **mean**, **count**, etc., to these groups.

pivot(): it is to create a pivot table by reshaping the data.

unstacked(): This is used to "unstack" (reshape) the data, turning the hierarchical index levels into columns.

```
Ввод [30]: group_type = df1.groupby('Type')['Price'].sum().sort_values(ascending=False)
```

```
Ввод [31]: group_type
```

```
Out[31]: Type
седан      3488840339
внедорожник 3005122999
кроссовер   2294975317
пикап       3026000000
минивэн     2468000000
универсал   2105400000
хэтчбек     1631000000
купе        1255000000
лифтбек     1044900000
фургон      5086999999
микроавтобус 1000000000
родстер     7900000000
кабриолет   2600000000
Name: Price, dtype: int64
```

```
Ввод [32]: group_car = df1.groupby('Brand')['Price'].sum().sort_values(ascending=False)
```

```
Ввод [33]: group_car
```

```
Out[33]: Brand
Toyota      2901515072
Hyundai      1353540703
Lexus        9099500000
Mercedes-Benz 6730200000
Kia          6580600000
BMW          4222570000
BA3          3484290000
Cadillac     3153000000
Chevrolet    3023899999
Nissan        2962105555
Volkswagen   2148250000
Li           1934000000
Dodge        1376500000
Infiniti     1291473288
Genesis      1190000000
Subaru       8420000000
Ford         7800000000
Mitsubishi   7700000000
Audi         7083000000
GA3          6999500000
Rolls-Royce   6300000000
Changan      5334000000
```

For the first, I used the **groupby()** function to find prices by the type and brand, so I can easily find the needed price.

```
Ввод [34]: grouped = df1.groupby('Color').agg(Total_Price=('Price', 'sum'), Avg_Price=('Price', 'mean')).sort_values(by='Total_Price', ascending=True)
```

```
Ввод [35]: grouped
```

```
Out[35]:
```

	Total_Price	Avg_Price
Color		
черный	3455206884	1.577720e+07
белый	3117709490	1.101664e+07
серый	917021282	8.817512e+06
серебристый	728639999	6.684771e+06
синий	330969999	6.364808e+06
металлик	276150000	1.022778e+07
коричневый	158100000	7.528571e+06
бордовый	113100000	7.540000e+06
жёлтый	103730000	1.728833e+07
зеленый	100462000	3.863923e+06
литые диски	97000000	1.616667e+07
красный	82680000	6.360000e+06
хамелеон	64500000	1.290000e+07
велюр	51100000	6.387500e+06
оранжевый	48100000	1.603333e+07
бежевый	39450000	5.635714e+06
вишня	39150000	5.592857e+06
голубой	33400000	5.566667e+06
Весь в родном окрасе	23499000	2.349900e+07

same thing for here, I used the **groupby()** function to group the data by car colors and applied the **agg()** function to calculate the total and average price for each color.

Color	001261 Автоцентр — это простота	080440 Автоцентр — это простота	129015 Автоцентр — это простота	184405 Автоцентр — это простота	3	Богатая комплектация Flagship	В наличии! Новая машина	Весь в родном окрасе	Двигатель 406	Дилерский центр «Chery Aстана»	... металлик	оранжевый
Brand												
Aston	0	0	0	0	0	0	0	0	0	0	...	0
Audi	0	0	0	0	0	0	0	0	0	0	...	6800000
BMW	0	0	0	0	0	0	0	0	0	0	...	65400000
BYD	0	0	0	0	0	0	0	0	0	0	...	0
Cadillac	0	0	0	0	0	0	0	0	0	0	...	0
Changan	0	0	0	0	0	0	0	0	0	0	...	0
Chery	0	0	0	0	0	0	0	0	0	7900000	...	0
Chevrolet	0	0	0	0	0	0	0	0	0	2890000	...	0
Chrysler	0	0	0	0	0	0	0	0	0	0	...	0
Daewoo	0	0	0	0	0	0	0	0	0	0	...	1550000
Dodge	0	0	0	0	0	0	0	0	0	0	...	0
DongFeng	0	0	0	0	0	0	0	0	0	0	...	0
EXEED	0	0	0	0	0	0	0	0	0	0	...	0
FAW	0	0	0	0	0	0	0	0	0	0	...	0
Ford	0	0	0	0	0	0	0	0	0	0	...	0
Geely	0	0	0	0	0	9700000	0	0	0	0	...	0
Genesis	0	0	0	0	15500000	0	0	0	0	0	...	0
Haval	0	0	0	0	0	0	0	0	0	0	...	0
Honda	0	0	0	0	0	0	0	0	0	0	...	0
Hyundai	0	7900000	11990000	7990000	0	0	0	0	0	0	...	70000000

The resulting **pivot table** will display the total prices (sum of the "Price" column) for each combination of car brand (rows) and car color (columns). For example, for the brand "**Audi**" and color "**металлик**" the total price will be the sum of all prices for **Audi** cars with that color. If no cars of a certain brand and color combination exist, the value will be 0.

[illegible]

Same, but I used **aggfunc** to calculate the total, average price, and count for each color in the pivot table

```
filtered_engine = df1[df1['Engine_Size'] > 2.5]
```

```
filtered_engine
```

	Title	Brand	Model	Year	Price	Condition	Type	Engine_Size	Fuel_Type	Transmission	Mileage	Color	Additional
1	Genesis G70	Genesis	G70	2021	23500000	Б/у	седан	3.3	бензин	КПП автомат	16500	белый	металлик
7	Toyota Highlander	Toyota	Highlander	2017	19500000	Б/у	кроссовер	3.5	бензин	КПП автомат	117000	черный	металлик, литые диски, тонировка, л...
13	Mercedes-Benz E 350	Mercedes-Benz	E 350	2013	12500000	Б/у	седан	3.5	бензин	КПП автомат	110000	литые диски	тонировка, панорамная крыша, хруст...
16	Toyota Camry	Toyota	Camry	2007	4000000	Б/у	седан	3.5	бензин	КПП автомат	413035	золотистый	литые диски, тонировка, кожа, аудио...
20	Infiniti QX80	Infiniti	QX80	2014	19000000	Б/у	внедорожник	5.6	бензин	КПП автомат	191000	бордовый	металлик, тонировка, люк, ветро...
...
933	Lexus LX 570	Lexus	LX 570	2017	46000000	Б/у	внедорожник	5.7	бензин	КПП автомат	70000	черный	металлик, Машина в идеальном состо...
934	Porsche Cayenne	Porsche	Cayenne	2006	8000000	Б/у	кроссовер	4.5	бензин	КПП автомат	155000	красный	металлик, литые диски, тонировка,...
938	Lexus LX 570	Lexus	LX 570	2016	40000000	Б/у	внедорожник	5.7	бензин	КПП автомат	270000	серебристый	металлик, тонировка, люк, сп...
940	Jeep Grand Cherokee	Jeep	Grand Cherokee	2014	14250000	Б/у	внедорожник	3.6	бензин	КПП автомат	164000	черный	металлик, литые диски, тонировка,...
942	Mercedes-Benz C 240	Mercedes-Benz	C 240	2000	3300000	Б/у	седан	2.6	бензин	КПП автомат	277000	металлик	тонировка, ксенон, велюр, аудиосистем...

Here I simply just filtered the data by creating a boolean condition where only rows with an "**Engine_Size**" greater than 2.5 are selected.

```
] grouped_c = df.groupby(['Brand', 'Color'])['Price'].sum()
```

```
] grouped_c
```

```
] Brand Color
Aston голубой 17000000
Audi вишня 3000000
голубой 3600000
зеленый 5850000
коричневый 3900000
...
УАЗ белый 9200000
жёлтый 4750000
зеленый 3000000
серебристый 2300000
серый 6599000
Name: Price, Length: 269, dtype: int64
```

```
1. grouped_c.groupby('Brand').agg({'Price': 'sum', 'Color': 'count'})
```

Here I want to calculate the total price for each combination of car brand and color, also using **groupby()** function.

```
Ввод [49]: unstacked = grouped_c.unstack(fill_value=0)
```

```
Ввод [50]: unstacked
```

```
Out[50]:
```

Color	001261 Автоцентр — это простота	080440 Автоцентр — это простота	129015 Автоцентр — это простота	184405 Автоцентр — это простота	3	Богатая комплектация Flagship	В наличии! Новая машина	Весь в родном окрасе	Двигатель 406	Дилерский центр «Chery Astana»	...	металлик	оранжевый
Brand													
Aston	0	0	0	0	0	0	0	0	0	0	...	0	0
Audi	0	0	0	0	0	0	0	0	0	0	...	6800000	0
BMW	0	0	0	0	0	0	0	0	0	0	...	65400000	0
BYD	0	0	0	0	0	0	0	0	0	0	...	0	0
Cadillac	0	0	0	0	0	0	0	0	0	0	...	0	0
Changan	0	0	0	0	0	0	0	0	0	0	...	0	0
Chery	0	0	0	0	0	0	0	0	0	7900000	...	0	0
Chevrolet	0	0	0	0	0	0	0	0	0	2890000	...	0	0
Chrysler	0	0	0	0	0	0	0	0	0	0	...	0	0
Daewoo	0	0	0	0	0	0	0	0	0	0	...	1550000	0
Dodge	0	0	0	0	0	0	0	0	0	0	...	0	0
DongFeng	0	0	0	0	0	0	0	0	0	0	...	0	0
EXEED	0	0	0	0	0	0	0	0	0	0	...	0	0
FAW	0	0	0	0	0	0	0	0	0	0	...	0	0
Ford	0	0	0	0	0	0	0	0	0	0	...	0	0
Geely	0	0	0	0	0	9700000	0	0	0	0	...	0	0
Genesis	0	0	0	0	15500000	0	0	0	0	0	...	0	0
Haval	0	0	0	0	0	0	0	0	0	0	...	0	0
Honda	0	0	0	0	0	0	0	0	0	0	...	0	0
Hyundai	0	7900000	11990000	7990000	0	0	0	0	0	0	...	70000000	0
Infiniti	0	0	0	0	0	0	0	0	0	0	...	0	0
JAC	0	0	0	0	0	0	0	0	0	0	...	0	0
Jaguar	0	0	0	0	0	0	0	0	0	0	...	0	0
Jeep	0	0	0	0	0	0	0	0	0	0	...	0	0

Then, we are coming to the **unstacked()** function, to view the aggregated data in a tabular format with car brands and colors. For example, if no "**Audi**" cars are available in the color "red," the corresponding value will be 0.

The result, as you can see, is a **DataFrame** with car brands as rows, car colors as columns, and the total price for each brand-color combination.

Data Analysis using other libraries such Patsy, Statsmodel and Scikit-learn

```
Ввод [60]: from patsy import dmatrices
```

Now, we are going to deal with the libraries that are also important for Data Analytics in Python.

1. **Patsy**: Create a formula or design matrices.
2. **Statsmodels**: Fit a statistical model to the data (e.g., linear regression, ANOVA).
3. **Scikit-learn**: Fit a machine learning model, evaluate it, and fine-tune hyperparameters.

```
Ввод [61]: y, X = dmatrices('Price ~ Year + Engine_Size + C(Type)', data=df1)

Ввод [62]: y
Out[62]: DesignMatrix with shape (952, 1)
Price
5950000
23500000
14790152
5200000
18000000
2550000
3950000
19500000
9100000
8700000
2500000
5350000
7000000
12500000
4800000
5200000
4000000
1500000

Ввод [63]: X
Out[63]: DesignMatrix with shape (952, 15)
Columns:
['Intercept',
 'C(Type)[Т. кабриолет]',
 'C(Type)[Т. кроссовер]',
 'C(Type)[Т. купе]',
 'C(Type)[Т. лифтбек]',
 'C(Type)[Т. микроавтобус]',
 'C(Type)[Т. минивэн]',
 'C(Type)[Т. пикап]',
 'C(Type)[Т. родстер]',
 'C(Type)[Т. седан]',
 'C(Type)[Т. универсал]',
```

It shows us the coefficients for **Year**, **Engine_Size**, and the dummy variables for **Type**. And we also get statistical details like p-values, R-squared, and other model diagnostics. This process helps to prepare data for statistical modeling. It automatically handles categorical variables, creates the necessary design matrices

Here we are attempting to perform a least-squares regression using the **Patsy** formula and **NumPy** for solving the linear system

Coefficients (coef): Tell us how much each factor (Year, Engine_Size, Type e.g.,) influences the price.

Residuals (resid): Tell us how far off the model's predictions are from the actual prices.

```
import statsmodels.api as sm
```

Here we are coming to statsmodel library

```
X2 = df1[['Year', 'Engine_Size']]  
y2 = df1['Price']
```

```
X2 = sm.add_constant(X2)
```

```
model = sm.OLS(y2, X2)
```

```
results = model.fit()
```

```
results.summary()
```

OLS Regression Results

Dep. Variable:	Price	R-squared:	0.595			
Model:	OLS	Adj. R-squared:	0.594			
Method:	Least Squares	F-statistic:	696.8			
Date:	Tue, 24 Dec 2024	Prob (F-statistic):	6.27e-187			
Time:	19:21:31	Log-Likelihood:	-16322.			
No. Observations:	952	AIC:	3.265e+04			
Df Residuals:	949	BIC:	3.266e+04			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-1.254e+09	5.32e+07	-23.581	0.000	-1.36e+09	-1.15e+09
Year	6.207e+05	2.64e+04	23.525	0.000	5.69e+05	6.72e+05
Engine_Size	5.762e+06	1.86e+05	30.938	0.000	5.4e+06	6.13e+06
Omnibus:	595.849	Durbin-Watson:	2.054			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	8703.824			
Skew:	2.616	Prob(JB):	0.00			
Kurtosis:	16.858	Cond. No.	4.88e+05			

We are using **Ordinary Least Squares (OLS) regression** using **statsmodels** to model the relationship between car prices (Price) and two independent variables:(Year and Engine_Size)

Here are the explanation of results of OLS Regression Results
Summary:

- **R-squared:** 0.595 (explains 59.5% of price variation).
- **Coefficients:**

Intercept: -1.254e+09 (not meaningful).

Year: 620,700 increase per year.

Engine_Size: 5.76 million increase per unit.

- **Significance:** Both Year and Engine_Size are statistically significant with very low p-values.
- **Diagnostics:** Residuals may not be normally distributed (Jarque-Bera and Omnibus tests), but no autocorrelation (Durbin-Watson = 2.054).

The model explains a good portion of price variation.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.88e+05. This might indicate that there are strong multicollinearity or other numerical problems.


```
Bвод [79]: df1['Price_above_10M'] = (df1['Price'] > 10000000).astype(int)
```

```
Bвод [80]: X3 = df1[['Year', 'Engine_Size']]
y3 = df1['Price_above_10M']
```

```
Bвод [81]: X3 = sm.add_constant(X3)
```

```
Bвод [82]: model = sm.Logit(y3, X3)
results = model.fit()
```

Optimization terminated successfully.
Current function value: 0.305531
Iterations 9

```
Bвод [83]: results.summary()
```

Out[83]: Logit Regression Results

Dep. Variable:	Price_above_10M	No. Observations:	952
Model:	Logit	Df Residuals:	949
Method:	MLE	Df Model:	2
Date:	Tue, 24 Dec 2024	Pseudo R-squ.:	0.5261
Time:	19:21:31	Log-Likelihood:	-290.87
converged:	True	LL-Null:	-613.73
Covariance Type:	nonrobust	LLR p-value:	6.058e-141
	coef	std err	z P> z [0.025 0.975]
const	-727.6010	56.947	-12.777 0.000 -839.214 -615.988
Year	0.3579	0.028	12.745 0.000 0.303 0.413
Engine_Size	2.1872	0.163	13.381 0.000 1.867 2.508

The model shows that both Year and Engine_Size significantly impact the likelihood of the car price being above 10 million. The **Pseudo R-squared** value of 0.5261 suggests a moderate fit. Both variables are highly significant predictors.

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import load_iris
```

Scikit-learn is one of the most popular libraries for machine learning in Python. It provides simple and efficient tools for data mining and data analysis.

```
: from sklearn.model_selection import train_test_split
: from sklearn.linear_model import LinearRegression
: from sklearn.metrics import mean_squared_error, r2_score
```

```
: X4 = df1[['Year', 'Engine_Size']]
: y4 = df1['Price']
```

```
:
:
: X_train, X_test, y_train, y_test = train_test_split(X4, y4, test_size=0.2, random_state=42)
:
: model = LinearRegression()
: model.fit(X_train, y_train)
:
: y_pred = model.predict(X_test)
:
: mse = mean_squared_error(y_test, y_pred)
: r2 = r2_score(y_test, y_pred)
: print(f"Mean Squared Error: {mse}")
: print(f"R-squared: {r2}")
```

```
Mean Squared Error: 38464954738395.836
R-squared: 0.6375018976123719
```

Linear Regression Results Summary:

Mean Squared Error (MSE): 38.46 trillion

Measures the average squared difference between actual and predicted prices. The high value suggests large prediction errors, likely due to the wide range of prices in the dataset.

R-squared (R^2): 0.638

Indicates that **63.8%** of the variance in car prices is explained by Year and Engine_Size.

Conclusion: The model provides a moderate fit, but the high **MSE** suggests further improvements are needed. Possible enhancements include scaling the data, adding more features, or using non-linear regression models.

```
: from sklearn.linear_model import LogisticRegression
: from sklearn.metrics import accuracy_score, confusion_matrix
```

```
: df1['Price_above_10M'] = (df1['Price'] > 10000000).astype(int)
:
: X5 = df1[['Year', 'Engine_Size']]
: y5 = df1['Price_above_10M']
:
: X_train, X_test, y_train, y_test = train_test_split(X5, y5, test_size=0.2, random_state=42)
:
: model = LogisticRegression()
: model.fit(X_train, y_train)
:
: y_pred = model.predict(X_test)
:
: accuracy = accuracy_score(y_test, y_pred)
: conf_matrix = confusion_matrix(y_test, y_pred)
:
: print(f"Accuracy: {accuracy}")
: print(f"Confusion Matrix:\n{conf_matrix}")
```

```
Accuracy: 0.7958115183246073
Confusion Matrix:
[[119  8]
 [ 31 33]]
```

1. **Accuracy (79.6%)** suggests reasonably good classification performance.
2. **Confusion Matrix** shows:

119 and 33 are correct predictions for each class.

8 and 31 are the misclassifications, where the model predicted incorrectly.

```
: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
from sklearn.datasets import load_iris

df1 = load_iris()
X6 = df1.data
y6 = df1.target

X_train, X_test, y_train, y_test = train_test_split(X6, y6, test_size=0.3, random_state=42)

rf = RandomForestClassifier(random_state=42)

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, n_jobs=-1, verbose=2)

grid_search.fit(X_train, y_train)

print("Best hyperparameters found: ", grid_search.best_params_)

best_rf = grid_search.best_estimator_
test_accuracy = best_rf.score(X_test, y_test)
print(f"Test set accuracy: {test_accuracy:.4f}")

Fitting 5 folds for each of 108 candidates, totalling 540 fits
Best hyperparameters found: {'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}
Test set accuracy: 1.0000
```

This code uses **GridSearchCV** to find the best *hyperparameters* for a Random Forest classifier on the Iris dataset. It splits the data into training and test sets (70/30), tests 108 parameter combinations with 5-fold cross-validation, and identifies the optimal settings: 100 estimators, no maximum depth, minimum split of 2, and minimum leaf of 1. The final model achieves 100% accuracy on the test set, indicating a perfect fit, which may suggest the dataset is simple or the model is overfitting.